
SimCert: A Classical-Simulability Audit of Trained Quantum-NLP Text Classifiers

Tanvir Ahmed

North South University, Dhaka, Bangladesh

tanvir.ahmed32@northsouth.edu

Abstract

Quantum natural language processing (QNLP) models, from compositional DisCoCat circuits to quantum self-attention networks and all-quantum token mixers, are promoted for a quantum benefit on text classification. That benefit is usually argued from accuracy, yet whether a *trained* QNLP circuit ever leaves the efficiently classically-simulable regime is rarely tested. We introduce SIMCERT, an open and model-agnostic protocol that audits this directly. It re-simulates each trained circuit as a bond-dimension- χ matrix product state, sweeps χ , and records the smallest value (written χ^*) at which a classical surrogate recovers the model’s own predictions, together with an entanglement-removal ablation and two dequantization witnesses. Auditing five published QNLP architectures across four datasets and up to ten qubits, we find that a classical surrogate recovers the predictions at small and non-growing bond dimension. Attention-style models are recovered at $\chi=1$, that is by a product state, and a compositional model at $\chi=2$; χ^* stays flat as the qubit count grows while the exact bound grows exponentially. The circuits do carry entanglement, but removing it barely moves the result. At the scales these models are tested, the quantum resource is therefore not the operative ingredient. This does not settle the asymptotic question, but it shifts the burden of proof: absent evidence that the required bond dimension grows with system size, the default reading of these models is classical. We release the audit as an open protocol and harness so that any QNLP model can be checked against this bar before a quantum benefit is claimed.

1 Introduction

Fault-tolerant quantum hardware does not yet exist, so most claims in quantum machine learning (QML) are established by classically simulating small circuits, and quantum NLP inherits the practice. A steady stream of QNLP models reports accuracy gains on text classification: compositional DisCoCat pipelines simulated for and run on real devices (Lorenz et al., 2021), quantum self-attention networks (Li et al., 2022; Chen et al., 2024), and, more recently, fully-quantum token mixers (Chen et al., 2025). The reported numbers are real. What is usually left implicit is the inference the reader is invited to draw, that the *quantum* structure is what does the work.

That inference is testable, and it is tested surprisingly rarely. Across the QNLP text-classification papers we examined (Appendix A), models are compared against classical *networks* of matched size but not against a classical *simulation of their own circuit*, which is the one baseline that would isolate whether the quantum resource, rather than the surrounding classical machinery or the trainable parameters, is responsible.¹ The same gap was recently closed for tabular QML (Bowles et al., 2024) and for quantum convolutional networks (Bermejo et al., 2024), resting on the theoretical footing that trainable variational models are constrained matrix product states (Shin et al., 2024) and that the absence of barren plateaus tends to imply classical simulability (Cerezo et al., 2025). So which QNLP results survive a classical surrogate of the trained circuit? That is the question we take up.

¹We count a paper as testing simulability only if it reports a classical surrogate of the trained circuit, meaning a tensor-network contraction, a Pauli-propagation estimate, or an explicit entanglement-removal ablation, rather than merely a separately-trained classical baseline. This is a focused review, not a saturated systematic survey; scope and counts are in Appendix A.

This paper makes three contributions, each an audit rather than a new model. First, we define a scoped simulability criterion for trained QNLP circuits, namely the smallest MPS bond dimension χ^* at which a classical surrogate recovers the model’s predictions, and we package it with an entanglement-removal ablation, the entanglement entropy, a geometric-difference kernel test, and a data-reuploading spectral test into a per-model simulability certificate. Second, we audit five published QNLP architectures, namely DisCoCat as realised in `lambeq`, the quantum self-attention models `QSANN` and `QMSAN`, the all-quantum token mixer `CLAQS`, and a variational-classifier reference, across four datasets, and we report how χ^* scales with system size. Third, we release an open, model-agnostic harness that consumes any trained circuit through a common intermediate representation, so the certificate can be recomputed and extended by others.²

The paper front-loads its epistemics. Section 2 places the work among the audits it extends. Section 3 sets out what a simulability audit can and cannot establish. Section 4 fixes definitions and the certificate. Sections 5 and 6 describe the audited models, the data, and the surrogate. Section 7 reports the findings, and Section 8 states the caveats and the open questions we think follow.

2 Related work

Our closest relatives are recent audits of quantum machine learning in other domains. Bowles et al. (2024) benchmark twelve QML models across many tabular datasets, match them against out-of-the-box classical baselines, and find that removing entanglement rarely hurts; we adopt their entanglement-removal ablation and their discipline of scoping every claim to the tested regime. Bermejo et al. (2024) show that trained quantum convolutional networks are effectively classically simulable by propagating low-weight observables, the vision-domain analogue of the question we ask for text. The theoretical footing is that a trained variational model is a constrained matrix product state (Shin et al., 2024), and that the very structure making such models trainable, the absence of barren plateaus, tends to make them classically simulable (Cerezo et al., 2025); our bond-dimension criterion is the operational instrument these results invite. Our two auxiliary witnesses come from the same program: the geometric difference between quantum and classical kernels (Huang et al., 2021) and the Fourier spectrum accessible to a data-reuploading encoding (Schuld et al., 2021).

On the QNLP side, the models we audit are the compositional DisCoCat framework (Lorenz et al., 2021), the Gaussian-projected and mixed-state quantum self-attention networks (Li et al., 2022; Chen et al., 2024), and the all-quantum token mixer (Chen et al., 2025). Each reports an accuracy comparison against a classical network, and none reports a classical surrogate of its own trained circuit, which is the gap we fill. What is new here is not the idea of a simulability audit but its first application to language, where the circuits are structurally unlike the tabular and vision cases (compositional cups, attention, and mixed states), and its packaging as a reusable protocol and harness rather than a one-off study. We read our results as extending the tabular and vision audits to language, and as agreeing with them, rather than as a challenge to any individual QNLP paper.

3 What a simulability audit can and cannot tell us

A skeptical result earns trust only if it anticipates how it might be wrong, so we state the limits before the evidence. An audit of this kind speaks to *trained* models at the scales they are actually run. It does not, and cannot, prove that no QNLP model at any size resists classical simulation, because the interesting asymptotic question is exactly the one simulation cannot reach. What it can do is check a concrete and load-bearing assumption, that the accuracy a model reports depends on quantum structure a cheap classical surrogate cannot reproduce.

Three failure modes shape the design. A surrogate that is too weak will fail to recover predictions for uninteresting numerical reasons, so we sweep the surrogate’s power, the bond dimension, rather than fixing it, and report the whole curve. Recovering accuracy and recovering the model’s function are different tests, since a surrogate could match the label while computing a different function, so we report agreement with

²Anonymized for review: <https://anonymous.4open.science/r/simcert>.




Figure 1: SIMCERT at a glance. A trained quantum-NLP classifier (top) embeds a sentence x , runs a quantum core that prepares $|\psi_\theta(x)\rangle$ and reads out Pauli expectations, and applies a classical head to obtain a label $\hat{y}$. SIMCERT replaces the quantum core with a classical matrix-product-state surrogate of bond dimension χ (bottom) and asks how small χ can be before the recovered label $\hat{y}^{(\chi)}$ changes. For every audited model the two tracks agree already at $\chi^* \leq 2$.




Figure 2: The nested hierarchy of classical simulability. A classically simulable model (CSIM) is a special case of one that is simulable up to a data-loading step (QESIM), which is a special case of quantum-simulable (QSIM) (Cerezo et al., 2025; Shin et al., 2024). SIMCERT certifies membership in CSIM whenever the trained state truncates to a small bond dimension, and we make no claim outside that innermost class.

the trained model’s own predictions and not only accuracy against ground truth. Finally, the absence of a quantum effect on a saturated benchmark says little, so we include a genuinely hard task and are explicit about where a benchmark is degenerate. None of this settles the asymptotic question. It settles the empirical one that current QNLP papers implicitly answer in the affirmative.

4 Framework and definitions

Setup. A QNLP text classifier maps an input x , a sentence or a sequence of token embeddings, to a trained circuit that prepares a state $|\psi_\theta(x)\rangle$ on n qubits, from which a readout, a set of Pauli expectation values possibly followed by a small classical head, produces a label $\hat{y}_\theta(x)$. We treat the trained parameters θ as fixed and ask about the function $\hat{y}_\theta$ (Figure 1).

Which notion of simulability. We use the standard hierarchy. A model is classically simulable (CSIM) if its outputs can be computed classically in polynomial time, and classically simulable up to a data-loading step, a weaker and common-in-practice notion, if the only obstruction is preparing the input state (Cerezo et al., 2025; Shin et al., 2024). Our surrogate is an MPS, which realises CSIM whenever the required bond dimension is small, and we make no claim about regimes where it is not. Every claim below has the form “simulable at bond dimension up to χ , at these model sizes, on these datasets” (Figure 2).

The bond-dimension criterion. For a trained state we form its optimal bond-dimension- χ MPS approximation $|\psi^{(\chi)}\rangle$ by sequential singular-value truncation, recompute the readout, and obtain a truncated prediction $\hat{y}^{(\chi)}(x)$ (Figure 4). Over the test set we report the truncated accuracy $A(\chi)$, the agreement $\pi(\chi) = \Pr_x[\hat{y}^{(\chi)}(x) = \hat{y}(x)]$ with the untruncated model, and the mean state fidelity $F(\chi)$. The certificate’s headline scalar is

$$\chi^* = \min\{\chi : A_{\text{full}} - A(\chi) \leq \tau\}, \quad (1)$$

Model	Paradigm	Cross-token mixing	Readout	State
disccocat	compositional	grammar cups (post-selection)	open-wire Born rule	pure
qsann	attention	Gaussian of measured Q/K (classical)	Pauli vector to head	pure
qmsan	attention	$\text{tr}(\rho_q \sigma_k)$ (mixed-state)	Pauli vector to head	mixed
claqs	token mixer	complex LCU and QSVT polynomial	XYZ Paulis to MLP	pure
vqc_text	reference	entangling ring	$\langle Z_0 \rangle$	pure

Table 1: The five audited models, one row per architecture, grouped by paradigm. **qmsan** is the only mixed-state model and is audited through the purification of its reduced density matrices. Full circuits, gate sets, and parameter counts are in Appendix B.

the smallest bond dimension whose accuracy is within a tolerance τ of the untruncated model. We set τ to the model’s own generalization gap, since differences smaller than the gap the model already exhibits are not meaningfully quantum resource, and we report robustness to two other tolerances in Appendix D. A model whose χ^* is small and does not grow with n is, in this operational sense, recovered by a cheap classical surrogate; we call such a model “bond-cheap” and use the term plainly thereafter.

Result (informal). *For every audited model at the tested scales, χ^* is a small constant, one or two, and does not grow when the qubit count is increased on a fixed task, while the bond dimension required for an exact MPS grows as $2^{\lfloor n/2 \rfloor}$.*

Corroborating witnesses. Bond dimension is one lens. The certificate adds three inexpensive ones. The first is an entanglement-removal ablation that retrain a product-state twin of the model and reports $\Delta A_{\text{ent}} = A_{\text{full}} - A_{\text{sep}}$, following Bowles et al. (2024). The second is the mean bipartite entanglement entropy $\bar{S}$ of the trained states, which lower-bounds the χ an exact representation needs. The third pairs two dequantization probes where they apply, the geometric difference g_{CQ} between the trained quantum kernel and a classical kernel (Huang et al., 2021) and the effective data-reuploading Fourier degree (Schuld et al., 2021). We report the full vector rather than collapsing to a single verdict, because, as Section 7 shows, the witnesses can legitimately disagree.

5 Models and datasets audited

Selection. We audit published QNLP text classifiers that target sentence or text classification, fit within a simulator’s reach at their published qubit counts, and can be reimplemented from their descriptions. Applying these criteria to the QNLP literature leaves a compact set spanning the two dominant paradigms, compositional and attention-based, plus a recent all-quantum mixer, and we add one plain variational classifier as a reference point. Table 1 compares them at a glance and stands in for a wall of per-model circuit diagrams; representative ansätze appear in Figure 3 and Appendix B. None of the attention or mixer models ship official code, so we reimplement from the equations and report the reproduction gap openly in Section 7, while the compositional model uses the official **lambeq** pipeline.

Datasets. We use the tasks these models were built on. MC, meaning classification of cooking versus computing, and RP, a relative-pronoun task, are the canonical DisCoCat benchmarks of Lorenz et al. (2021). MC is small and near-separable, whereas RP is the harder, low-data task where published accuracies actually differ across models. We add a balanced subset of SST-2 as a genuinely non-separable sentiment task, and a synthetic MC-style set for controlled scaling. Sizes and preprocessing are in Appendix C. Two limitations are worth stating here rather than burying. Our RP reimplementations use our own token embeddings, and our DisCoCat pipeline uses an offline reader in place of the trained grammar parser whose model host is no longer reachable. Both cost accuracy on RP, as we report.




Figure 3: A representative variational ansatz (the reference classifier; the attention and mixer models use the same primitives). Each input feature is angle-encoded ($R_y(x_i)$, orange), followed by a trainable layer ($R_y(\theta_i)$, green), a CNOT entangling chain, and a trainable $R_z(\phi_i)$. The block repeats L times with the data re-uploaded each layer, and the class is read from $\langle Z_0 \rangle$. The audit forms the bond-dimension- χ MPS approximation of the state prepared *just before* measurement. Per-model circuits are in Appendix B.

Algorithm 1 SIMCERT bond-dimension audit for one trained model

```

1: input: trained model  $M$ , test set  $D$ , tolerance  $\tau$ , bond dimensions  $X$ 
2: for each example  $x \in D$  do
3:   export the trained units of  $x$  (one whole-sentence circuit, or per-token circuits)
4:   for each  $\chi \in X$  do
5:     truncate each unit's state to bond dimension  $\chi$ ; recompute its Pauli readout
6:      $\hat{y}^{(\chi)}(x) \leftarrow$  model's own decision rule on the truncated readout
7:   end for
8: end for
9:  $A(\chi), \pi(\chi), F(\chi) \leftarrow$  accuracy, agreement, fidelity aggregated over  $D$ 
10:  $\chi^* \leftarrow \min\{\chi : A_{\text{full}} - A(\chi) \leq \tau\}$  ▷ Eq. 1
11: return certificate  $(\chi^*, \Delta A_{\text{ent}}, \bar{S}, g_{CQ}, \text{Fourier degree})$ 

```

6 The classical surrogate and simulation protocol

We do not conjecture simulability, we realise it. Each trained model exports its per-example circuit to a common representation, OpenQASM for gate circuits, or, for the token mixer whose state is a matrix polynomial rather than a gate sequence, the trained statevector directly. A single audit consumer then re-simulates every model identically (Algorithm 1), so the audit never depends on which framework produced a circuit. The MPS truncation itself is an exact sequential-SVD computation, cross-checked against an independent tensor-network simulator. On circuits with known answers, a GHZ state and a product state, it reproduces the analytic bond structure, which is our correctness check for the whole pipeline.

The surrogate's cost is explicit. An MPS sweep on an n -qubit, depth- d circuit scales as $\mathcal{O}(nd\chi^2)$, and the exact-representation bound is $\chi = 2^{\lceil n/2 \rceil}$, so a small measured χ^* against a rapidly growing exact bound is precisely the signal we are after. All experiments run on a single CPU with no GPU. Where a reimplemention is heavier than the hardware comfortably allows we say so and reduce scale rather than quietly dropping runs, and full configurations are in Appendix D.

7 Results

We report three findings, then a subsection of effects we looked for and did not see. Unless noted, numbers are the mean over three seeds; the aggregate certificate is Table 2 and the reproduction gap is Table 3.

7.1 A classical surrogate recovers the predictions at small bond dimension

Across models and datasets, a bond-dimension-one or bond-dimension-two surrogate recovers the trained model's predictions (Figure 6). On the canonical MC task the attention and mixer models **qsann**, **qmsan**, and




Figure 4: Following the dequantization picture of Shin et al. (2024), the trained model’s expectation value is a *coefficient* matrix product state $\mathbf{C}^q(\theta)$ (blue) contracted with a fixed set of basis functions $\mathbf{T}(x)$ (red), and its bond dimension caps the entanglement it can carry. If that bond dimension is small, a classical MPS of the same size (teal) contracts with the same basis functions to reproduce the output, so the model is basis-equivalent to a cheap classical one. SIMCERT measures this bond dimension directly: it truncates the trained state to bond dimension χ and finds the smallest χ that still recovers the predictions, which for every audited model is 1 or 2, against an exact-representation bound of $2^{\lfloor n/2 \rfloor}$.

claqs, together with the reference classifier, all reach full accuracy at $\chi=1$, a product state, with no accuracy lost; their curves in Figure 6 are flat down to the smallest bond dimension. The compositional **discocat** is the exception and needs $\chi=2$, but no more. The pattern holds on real data. Our DisCoCat run reproduces the published MC accuracy to within 0.02, at 0.778 against 0.798 (Table 3), and is recovered at $\chi=2$. The scales are already low, so in the operational sense of Section 4 these trained models are bond-cheap.

7.2 χ^* does not grow with system size

The obvious objection is that small circuits are trivially simulable, so the audit only bites at scale. Figure 5 tests this directly. On a fixed task we increase the qubit count of the reference model from four to ten and track χ^* . It stays flat at 1 while the exact bond bound grows from 4 to 8 to 16 to 32. Adding qubits neither raises the bond dimension a classical surrogate needs nor, on this task, improves accuracy, which drifts from 0.63 at four qubits to 0.52 at ten. *The gap between what the model uses and what it could use widens with every qubit.*

7.3 Removing the quantum resource barely moves the result

Two further witnesses agree (Figure 7). The trained circuits are not trivial product states, since their mean bipartite entanglement entropy runs from about 0.2 to 1.0 nats on MC (Figure 7a) and from about 0.1 to 1.3 nats across all datasets (Table 2), yet truncating to $\chi=1$ loses no predictions for the attention models (Figure 7b), so the entanglement they build is not load-bearing for their decisions. The entanglement-removal ablation tells the same story on the easy tasks, where a retrained product-state twin matches the full model and ΔA_{ent} is within noise of zero for **qsann** and **qmsan** on MC. The one place the state fidelity climbs gradually with χ is the eight-qubit mixer **claq**s, where fidelity rises from 0.72 at $\chi=1$ to 0.99 at $\chi=4$; this is also the one place the bond-dimension axis has room to resolve structure, and even there the predictions are recovered at $\chi=1$.

7.4 Effects we did *not* observe

We looked for, and did not find, three things. We did not see χ^* grow with qubit count on a fixed task (Figure 5). We did not see the entanglement-removal ablation cost accuracy on the tasks where the models saturate. And we did not see the compositional model behave like the attention models. **discocat** alone




Figure 5: The required bond dimension χ^* does not grow with qubit count. Left, accuracy retained versus bond dimension for the reference model at $n \in \{4, 6, 8, 10\}$ on SST-2; each curve is flat down to $\chi=1$. Right, the measured χ^* (orange) stays at 1 while the exact-MPS bound $2^{n/2}$ (grey, dashed) grows exponentially on the log axis. A product-state surrogate recovers the predictions at every size tested. Accuracy here is modest and does not rise with n , so this probes recovery rather than a strong classifier.




Figure 6: Two views of the same bond-dimension sweep, one curve per model and dataset, averaged over seeds. Color encodes the model and marker shape encodes the dataset, so each series stays distinguishable in grayscale. Panel (a) plots test accuracy retained against χ . A curve that is flat down to $\chi=1$ is the signature of a bond-cheap decision, since the classical surrogate loses no predictions as it is weakened to a product state. The only curves that climb are the three compositional **discocat** runs, which recover their accuracy at $\chi=2$; every other series is flat. Panel (b) plots the state fidelity $F(\chi)$ between the truncated and the full circuit for the same runs. This climbs from as low as 0.2 at $\chi=1$ toward one, so the trained states do carry entanglement. Read together the panels reconcile the finding: the wavefunctions are entangled, yet the decisions built on them do not need that entanglement at the sizes tested.

needs $\chi=2$, and truncating it to $\chi=1$ collapses its accuracy toward chance, so its grammar cups create entanglement that is, minimally, load-bearing. This lone exception matters more than the uniform cases, because it shows the audit is not simply reporting $\chi=1$ by construction.




Figure 7: Entanglement is present but, with one exception, not load-bearing, on MC for all five models). Left, the mean bipartite entanglement entropy $\bar{S}$ is non-zero for every model, so the trained states are genuinely entangled. Right, the full model and its $\chi=1$ product-state surrogate reach the same accuracy for the attention and mixer models, whereas for **discocat** the $\chi=1$ surrogate drops toward chance, marking the one place where the entanglement built by its grammar cups is load-bearing.

Model	Data	seeds	Full acc	Acc@ $\chi=1$	χ^*	$\bar{S}$ (nats)	ΔA_{ent}
claq	mc	3	1.000±0.000	1.000±0.000	1	1.017±0.158	–
claq	sst2	1	0.520	0.525	1	0.517	–
discocat	mc	3	0.936±0.065	0.551±0.079	2	0.659±0.013	–
discocat	mc_real	3	0.778±0.087	0.500±0.000	2	0.667±0.019	–
discocat	rp	3	0.505±0.133	0.387±0.000	1/2	0.654±0.005	–
qmsan	mc	3	1.000±0.000	1.000±0.000	1	0.169±0.077	0.000±0.000
qmsan	rp	3	0.581±0.115	0.591±0.110	1	0.106±0.022	-0.065±0.053
qsann	mc	3	1.000±0.000	1.000±0.000	1	0.295±0.064	0.000±0.000
qsann	rp	3	0.505±0.092	0.548±0.070	1	0.400±0.012	-0.022±0.030
vqc_text	mc	3	1.000±0.000	1.000±0.000	1	0.836±0.115	0.051±0.036
vqc_text	rp	3	0.634±0.080	0.634±0.080	1	0.685±0.059	0.097±0.121
vqc_text	sst2	7	0.610±0.050	0.610±0.050	1	1.282±0.232	0.123±0.033

Table 2: Simulability certificate for all models. Mean and standard deviation over seeds. Acc@ $\chi=1$ is the accuracy of a product-state surrogate; χ^* is Eq. 1; $\bar{S}$ is the mean bipartite entanglement entropy in nats; ΔA_{ent} is the entanglement-removal gap, with a dash where a separable twin is not defined. Attention and mixer models are recovered at $\chi=1$, and **discocat** at $\chi=2$.

Reproduction, honestly. Table 3 and Figure 8 place our reimplementations against the published numbers. On MC the compositional model reproduces to within 0.02. On RP, a task of only 74 training items that the original papers themselves flag as fragile, our numbers run about 0.17 to 0.22 below published, which we attribute to our own token embeddings and, for DisCoCat, the offline reader standing in for the unavailable grammar parser. We flag this plainly, because it matters for one thing and not another. It weakens any claim about matching RP accuracy, but the simulability verdict is about our own trained circuits and is unaffected by whether they hit the published accuracy. The certificate audits the model in hand. One might still object that better-trained models would demand more bond dimension, so that our gap, rather than any real property, drives the verdict. The evidence cuts the other way. On MC, where we reproduce the published accuracy, the attention models are already recovered at $\chi=1$, so the $\chi=1$ finding does not depend on the RP shortfall; and χ^* is flat in system size on a task we do fit (Figure 5), which is the opposite of what a training-quality artefact would produce.

Model	Data	seeds	ours (test)	published
discocat	mc_real	3	0.778 ± 0.087	0.798
discocat	rp	3	0.505 ± 0.133	0.723
qmsan	rp	3	0.581 ± 0.115	0.756
qsann	rp	3	0.505 ± 0.092	0.677
vqc_text	sst2	7	0.610 ± 0.050	–

Table 3: Reproduction gap: our audited accuracy against the published number, on the datasets with a published value. MC_REAL reproduces to within 0.02, whereas RP runs below published, being a fragile task of only 74 training items reimplemented with our own embeddings and, for DisCoCat, an offline reader. The simulability verdict concerns our own trained circuits and is independent of this gap.




Figure 8: Reproduction gap: our audited accuracy against the published number for the pairs with a published value; the annotation is ours minus published. MC_REAL reproduces closely, whereas RP, a fragile task of only 74 training items reimplemented with our own embeddings and an offline reader, runs below. The simulability verdict concerns our own trained circuits and does not depend on closing this gap.

8 Discussion, caveats, and open questions

The scoped claim is narrow and, we think, well supported. The QNLP models we audited reproduce their behavior under a classical surrogate of small and non-growing bond dimension, and the entanglement they carry is, with one revealing exception, not load-bearing for their predictions, at the model sizes and datasets tested. That scope is a real limit rather than a hedge, since small-scale simulation cannot reach the asymptotic question, but what it does reach it answers cleanly.

Several caveats bound the reading. The audit is average-case over a test set, not a statement about every input. χ^* is defined against a tolerance, and while we check robustness to its choice, a different tolerance shifts the borderline cases. Our simulability notion is the practical one, up to data loading, not a hardness result. The scales are small, and although the point of the scaling study is that they need not be large for the pattern to appear, a determined optimum at larger n remains out of simulation’s reach. Our reimplementations also fall short of the published accuracies on RP, so for those models we audit weaker instances than the originals, though the verdict concerns the circuits we actually trained rather than the published ones. Finally, the all-quantum mixer runs at reduced scale on a CPU, where its accuracy is near chance, so we treat it as a demonstration that the audit works at eight qubits rather than as an audit of the published model. That model remains untested at its full scale, which would need a GPU to reproduce, and we say so plainly rather than presenting a reduced-scale run as a full audit.³

³This is the only place additional compute would change what we can measure; the rest of the audit is CPU-bound by design.

There is a constructive reading. If these models are bond-cheap at deployment scale, their quantum implementation earns its overhead only where a classical MPS cannot follow, that is where χ^* grows with system size, and none of the audited models show it growing. The audit therefore names the target a genuine quantum ingredient must hit rather than merely cataloguing a shortfall. That reframing turns the negative headline into a design target and suggests concrete questions for QNLP. Which architectural choices, if any, make χ^* grow with n ? Do the compositional cups that make **discocat** need $\chi=2$ point toward structure worth scaling? Should QNLP benchmarks report a simulability certificate alongside accuracy, the way tabular QML increasingly does? Our position is collegial rather than combative, since the tabular and vision audits reached the same conclusion in their domains (Bowles et al., 2024; Bermejo et al., 2024), and we read our results as extending, not contradicting, theirs. *Until a QNLP model shows its bond dimension growing, the burden of proof for a quantum benefit sits with the model, not with the skeptic.*

Reproducibility statement

The harness, all model reimplementations, exact configurations, seeds, dataset indices, and the scripts that regenerate every figure and table from stored results are released at the anonymized repository above. The audit is deterministic given a trained model, and the only stochasticity is in training, which we control by seed.

References

- J. Bowles, S. Ahmed, and M. Schuld. Better than classical? The subtle art of benchmarking quantum machine learning models. *arXiv:2403.07059*, 2024.
- P. Bermejo et al. Quantum convolutional neural networks are (effectively) classically simulable. *arXiv:2408.12739*, 2024.
- M. Cerezo et al. Does provable absence of barren plateaus imply classical simulability? *Nature Communications*, 2025. *arXiv:2312.09121*.
- S. Shin, Y. S. Teo, and H. Jeong. Dequantizing quantum machine learning models using tensor networks. *Phys. Rev. Research* 6:023218, 2024.
- H.-Y. Huang et al. Power of data in quantum machine learning. *Nature Communications* 12:2631, 2021. *arXiv:2011.01938*.
- M. Schuld, R. Sweke, and J. J. Meyer. Effect of data encoding on the expressive power of variational quantum machine learning models. *Phys. Rev. A* 103:032430, 2021.
- R. Lorenz et al. QNLP in practice: running compositional models of meaning on a quantum computer. *arXiv:2102.12846*, 2021.
- G. Li, X. Zhao, and X. Wang. Quantum self-attention neural networks for text classification. *arXiv:2205.05625*, 2022.
- F. Chen et al. Quantum mixed-state self-attention network. *arXiv:2403.02871*, 2024.
- J. Chen et al. CLAQS: compact learnable all-quantum token mixer for text classification. *arXiv:2510.06532*, 2025.

A Literature-review methodology

The motivating claim in Section 1, that QNLP text-classification papers argue a quantum benefit from accuracy but rarely test whether the *trained circuit* is classically simulable, rests on a focused review rather than an exhaustive systematic survey, and we scope it honestly here.

Search and inclusion. We searched arXiv (quant-ph, cs.CL) and Semantic Scholar for the conjunctions of {"quantum"} with {"natural language processing", "text classification", "self-attention", "sentence classification"} over 2019-2025, and followed citations from the models we reimplement. We include a paper if it (i) proposes a quantum or hybrid model for a text-classification task and (ii) reports an accuracy comparison. We record, for each, whether it reports a *classical surrogate of the trained circuit*, meaning a tensor-network contraction, a Pauli-propagation estimate, or an explicit entanglement-removal ablation, as opposed to merely a separately-trained classical network of comparable size.

What we found. Table 4 lists the core set. All report a classical *network* baseline; none report a classical surrogate of their own trained circuit, and none report the bond dimension, entanglement, or a simulability check of the trained model. This is the gap SIMCERT fills. We stress the limits: this is a focused review of the papers central to this study plus those surfaced by the searches above, not a saturated systematic survey, and we did not attempt inter-rater agreement. We therefore phrase the claim in the main text as scoped to "the papers we examined," and treat a full systematic survey (with pre-registered search strings and dual coding, in the style of the tabular-QML audit of Bowles et al., 2024) as future work.

Paper	classical-network baseline	trained-circuit surrogate	reports χ / entanglement
DisCoCat / lambeq (Lorenz et al., 2021)	- ^a	no	no
QSANN (Li et al., 2022)	yes	no	no
QMSAN (Chen et al., 2024)	yes	no	no
CLAQS (Chen et al., 2025)	yes	no	no

Table 4: Core reviewed set. ^aDisCoCat compares to a bag-of-words/word-sequence baseline rather than a matched network. None reports a classical surrogate of the trained circuit or a simulability diagnostic.

B Model circuits and glossary

All models are reimplemented in PennyLane; lambeq produces the DisCoCat circuits. We give each model’s trained state below, and Figure 9 shows a representative circuit for each; the audit truncates that state (Algorithm 1).

vcq_text (reference). Bag-of-embeddings feature $x \in \mathbb{R}^n$ angle-encoded on n qubits, L layers of $R_y(x_i)$ (re-uploaded), trainable R_y, R_z , and a CNOT ring; readout $\langle Z_0 \rangle$ (Fig. 3).

QSANN (Gaussian-projected quantum self-attention). Per token, the same strongly-entangling template ($H^{\otimes n}$, then R_x/R_y layers with CNOT chains) prepares $|\psi_s\rangle = U_{enc}(x_s)H^{\otimes n}|0\rangle$; three trainable ansätze U_q, U_k, U_v give a query scalar $\langle Z_0 \rangle$, a key scalar, and a d -dimensional value vector of single-qubit Pauli expectations. Attention $\alpha_{s,j} = \exp(-(\langle Z_q \rangle_s - \langle Z_k \rangle_j)^2)$, the residual sum, mean-pool, and sigmoid are all classical. The audit therefore acts *per token* on the pure states $U_{q,k,v}|\psi_s\rangle$, and the classical head is the composition function.

QMSAN (mixed-state self-attention). Per token, a data-reuploading circuit ($R_x(x_i)$; $L \times$ [IsingZZ ring, R_y layer]; re-upload) prepares a pure state whose reduced density matrix on the first $n/2$ qubits is the query/key object; attention is the Hilbert-Schmidt overlap $\alpha_{s,j} = \text{tr}(\rho_{q,s}\sigma_{k,j})$. We audit this on the *purification*: we truncate the pure n -qubit state to bond dimension χ , then take its partial trace to form ρ, σ . Values are per-qubit $\langle Z_i \rangle$.

CLAQS (all-quantum token mixer). On $q=8$ data qubits, each token’s ansatz-14 embedding unitary U_j is combined by a learnable, ℓ_1 -normalised complex linear-combination-of-unitaries $M = \sum_j \bar{b}_j U_j$; a learnable QSVT polynomial $P_c(M) = \sum_{k=0}^5 c_k M^k$ and a window feed-forward unitary U_{FF} produce $|\psi\rangle = \text{normalise}(U_{FF} P_c(M) |0^8\rangle)$; readout is 24 XYZ Pauli expectations into a two-layer MLP. Because $|\psi\rangle$ is a matrix polynomial rather than a gate sequence, the audit consumes the trained statevector directly. At $q=8$ the exact MPS bond dimension is 16, which is why this is the one model where the χ -axis has room to resolve structure.




Figure 9: Representative circuits for the four audited models (simplified; exact per-model circuits are released with the harness). Gates are colored by role. **QMSAN** is read out as a reduced density matrix, hence the mixed-state audit path; **DisCoCat** post-selects the cup qubits and reads the single open wire by the Born rule.

discoCat. A sentence’s pregroup structure becomes one circuit: single-qubit word states (R_x / Euler $R_x R_z R_x$), IQP ansätze for multi-qubit words (per layer: H on all wires, then a chain of adjacent CR_z), a parameter-free GHZ for the relative pronoun, and grammar cups realised as Bell effects (post-select the paired qubits on $|0\rangle$). We audit the whole-sentence pure state before post-selection, then post-select and read the open wire’s Born rule.

Glossary (for NLP readers). **NISQ**: noisy intermediate-scale quantum hardware. **PQC**: parameterised quantum circuit. **Ansatz**: a fixed circuit template whose rotation angles are trainable. R_x, R_y, R_z : single-qubit rotations; **CNOT**/ CR_z : two-qubit entangling gates. **Statevector**: the 2^n complex amplitudes of an n -qubit pure state. **MPS** (matrix product state): a tensor-network factorisation whose **bond dimension** χ caps the entanglement it can represent; $\chi=1$ is a product (unentangled) state. **Entanglement entropy** $\tilde{S}$: bipartite von Neumann entropy of the state (nats). **IQP** / **LCU** / **QSVT**: instantaneous-quantum-polynomial ansatz / linear combination of unitaries / quantum singular-value transformation. **CSIM**/**QESIM**: classically simulable / classically simulable up to data loading.

C Datasets and preprocessing

MC (meaning classification). We use two variants. MC-REAL is the Lorenz et al. set (70 train / 30 dev / 30 test; 17-word vocabulary; 3-4 word cooking-vs-computing sentences), obtained from the authors’ public resources repository; POS tags are stripped (**word_N→word**) since our bag-of-token models do not use them and the DisCoCat reader re-parses. MC-SYNTH is a self-contained stand-in generated from disjoint food/IT subject-verb-object templates (65 per class), used only for the controlled scaling study; being disjoint-vocabulary it is near-linearly-separable, which is why every model saturates on it.

RP (relative pronoun). The Lorenz et al. RELPRON-derived set (74 train / 31 test; 115-word vocabulary; four-token noun phrases), distinguishing subject- from object-relative clauses. We carve a 20% validation split from train. RP is the discriminating benchmark and, with only 74 training items, the fragile one, and the original papers note near-degenerate behaviour here, which we reproduce.

SST-2. A balanced subset of the Stanford Sentiment Treebank binary task (600 train / 200 test, drawn from the train and labelled-validation splits respectively via the HuggingFace datasets server), sequence length capped at 16 tokens. A genuinely non-separable, real-text task; accuracies here are in the 0.5-0.65 range, not saturated.

Embeddings and tokenisation. For `qsann`, `qmsan`, `claqs`, and `vqc_text` we use a trainable per-token embedding (dimension equal to the model’s feature width); `vqc_text` averages token embeddings (bag-of-embeddings), the attention models keep per-token vectors, and `claqs` maps each token embedding through a trainable linear layer to ansatz-14 angles. DisCoCat does not use embeddings: each word’s parameters *are* its gate angles, trained via the `lambeq` pipeline with the offline `cups_reader` (the trained Bobcat grammar parser’s model host is no longer reachable, an acknowledged fidelity cost on RP). Splits are seed-stable; the exact indices and the synthetic generator are released with the harness so every split is reproducible.

D Configurations, compute, and robustness

Hyperparameters. Table 5 lists the per-model settings. Quantum models train with Adam (AdamW for `claqs`); `discocat` trains with SPSSA ($a=0.05, c=0.06$) as in the original pipeline. Quantum-model weights are initialised $\mathcal{N}(0, 0.01)$; all runs use seeds $\{1, 2, 3\}$. The matched classical baseline is logistic regression on bag-of-words counts. Quantum-model feature widths follow the source papers ($n=2$ qubits on MC, $n=4$ on RP; $q=8$ for `claqs`).

Model	qubits	depth/layers	optimiser	epochs
<code>vqc_text</code>	4-10	2	Adam (0.05)	30-40
<code>qsann</code>	2 / 4	$D_{enc}=1, D_{qkv}=1$	Adam (0.05)	40
<code>qmsan</code>	2 / 4	$L=1$ (IsingZZ ring)	Adam (0.05)	40
<code>claqs</code>	8	$L=1$, QSVT degree 5	AdamW (0.05)	10-12
<code>discocat</code>	3-7	IQP $L=1$	SPSSA	120

Table 5: Per-model training configuration. Full configs and the exact commands are released with the harness.

Compute. All experiments run on a single Apple M1 CPU (8 GB RAM, no GPU). The audit is deterministic given a trained model, since expectation values are computed exactly rather than sampled, so the only stochasticity is training. Statevector simulation is memory-light at these sizes (a 16-qubit state is ≈ 1 MB); the binding constraint is training the attention models, not the audit. `claqs` is the exception: a faithful full-scale reproduction of its published 91.6% on SST-2 needs a GPU, so we run it at reduced scale (100 training examples, ≤ 6 -token windows, 10 epochs); its SST-2 accuracy is consequently near chance and we report the audit, not an accuracy reproduction.

Robustness of χ^* to the tolerance. Equation 1 defines χ^* against a tolerance τ . We compute it under three definitions: τ_{gen} (the model’s train-test generalization gap, our default), τ_{ci} (the half-width of a 95% bootstrap CI on full- χ accuracy), and τ_{agree} (smallest χ with prediction agreement $\pi(\chi) \geq 0.95$). Across all runs the three agree on χ^* in the large majority of cases, namely $\chi=1$ for the attention and mixer models and $\chi=2$ for `discocat`, and where they differ they do so by a single bond-dimension step on borderline, low-accuracy RP runs. The headline conclusion is therefore not an artefact of the tolerance choice.